

In [38]:

```
#import necessary libraries
import pandas as pd
import numpy as np
import scipy as sp
import matplotlib.pyplot as plt
import seaborn as sns

#read data.csv into a dataframe df and check basic information
df=pd.read_csv('/Users/danhongcheng/desktop/data.csv')
df.describe()
```

/Users/danhongcheng/opt/anaconda3/lib/python3.8/site-packages/IPython/core/interactiveshell.py:3165: DtypeWarning: Columns (18,56) have mixed types.Specify dtype option on import or set low_memory=False.

has_raised = await self.run_ast_nodes(code_ast.body, cell_name,

Out[38]:

	id	member_id	loan_amnt	funded_amnt	funded_amnt_inv	int_rate	installment	annual_inc	dt
count	1.999990e+05	0.0	199999.000000	199999.000000	199999.000000	199999.000000	199999.000000	1.999990e+05	199997.000000
mean	6.229635e+07	NaN	15278.117141	15278.117141	15269.472943	12.361755	441.383672	7.815068e+04	19.164496
std	3.941126e+06	NaN	8651.138790	8651.138790	8646.324513	4.242108	247.052659	8.051380e+04	9.157591
min	5.670500e+04	NaN	1000.000000	1000.000000	900.000000	5.320000	14.770000	0.000000e+00	0.000000
25%	5.941173e+07	NaN	8500.000000	8500.000000	8475.000000	9.170000	261.880000	4.757100e+04	12.540000
50%	6.221754e+07	NaN	14000.000000	14000.000000	14000.000000	12.290000	383.810000	6.500000e+04	18.610000
75%	6.564457e+07	NaN	20000.000000	20000.000000	20000.000000	14.650000	580.730000	9.340000e+04	25.410000
max	6.861706e+07	NaN	35000.000000	35000.000000	35000.000000	28.990000	1445.460000	9.000000e+06	999.000000

8 rows × 115 columns

In [39]:

```
df.shape
```

Out[39]: (199999, 148)

In [40]:

```
df.sample(10)
```

Out[40]:

	id	member_id	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	grade	sub_grade	...	hardship_payoff_l
611	68606692	NaN	15000	15000	15000.0	36 months	12.88	504.55	C	C2	...	
118631	61480137	NaN	28000	28000	28000.0	36 months	7.26	867.90	A	A4	...	
145656	60185273	NaN	8000	8000	8000.0	36 months	10.99	261.88	B	B4	...	
75555	64097871	NaN	35000	35000	35000.0	60 months	16.99	869.66	D	D3	...	
16530	67485341	NaN	27600	27600	27600.0	60 months	10.64	595.15	B	B4	...	
184121	57135502	NaN	10000	10000	10000.0	36 months	12.69	335.45	C	C2	...	
76528	63179137	NaN	18000	18000	18000.0	36 months	10.99	589.22	B	B4	...	
121651	60785001	NaN	4000	4000	4000.0	36 months	7.89	125.15	A	A5	...	
182855	57296281	NaN	28000	28000	28000.0	60 months	9.99	594.78	B	B3	...	
148176	59709156	NaN	35000	35000	35000.0	36 months	12.69	1174.07	C	C2	...	

10 rows × 148 columns

In [41]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 199999 entries, 0 to 199998
Columns: 148 entries, id to settlement_term
dtypes: float64(70), int64(45), object(33)
memory usage: 225.8+ MB
```

In [42]:

```
df['loan_status'].value_counts(dropna=False)
```

Out[42]: Fully Paid 140991

```
Charged Off      35090
Current          22637
Late (31-120 days)  785
In Grace Period   347
Late (16-30 days)  148
Default           1
Name: loan_status, dtype: int64
```


```
In [43]: # slice the dataframe so as to only keep records that have 'loan_status' in ['Fully Paid', 'Charged off']
df=df[df['loan_status'].isin(['Fully Paid', 'Charged Off'])]
print(df['loan_status'].value_counts())
df.shape
```

```
Fully Paid      140991
Charged Off      35090
Name: loan_status, dtype: int64
```

```
Out[43]: (176081, 148)
```

```
In [44]: #caculate missing rates of each feature and plot histogram of feature missing rates
missing_rate=df.isnull().mean().sort_values(ascending=False)
plt.hist(x=missing_rate, bins=20)
plt.xticks(np.arange(0,1,0.05), rotation='vertical')
plt.title('Histogram of Feature Missing Rates')
plt.xlabel('Rates of Missing Data')
plt.ylabel('Count')
```

```
Out[44]: Text(0, 0.5, 'Count')
```



In [45]:

```
#by visualizing the missing rate of features, we can set the threshold of missing rate to 0.2 to drop features
drop_feature=sorted(list(missing_rate[missing_rate.values>0.05].index))
print(drop_feature)
```

```
['all_util', 'annual_inc_joint', 'debt_settlement_flag_date', 'deferral_term', 'desc', 'dti_joint', 'emp_length', 'emp_title', 'hardship_amount', 'hardship_dpd', 'hardship_end_date', 'hardship_last_payment_amount', 'hardship_length', 'hardship_loan_status', 'hardship_payoff_balance_amount', 'hardship_reason', 'hardship_start_date', 'hardship_status', 'hardship_type', 'il_util', 'inq-fi', 'inq_last_12m', 'max_bal_bc', 'member_id', 'mths_since_last_delinq', 'mths_since_last_major_derog', 'mths_since_last_record', 'mths_since_rcnt_il', 'mths_since_recent_bc_dlq', 'mths_since_recent_inq', 'mths_since_recent_revol_delinq', 'next_pymnt_d', 'num_tl_120dpd_2m', 'open_acc_6m', 'open_act_il', 'open_il_12m', 'open_il_24m', 'open_rv_12m', 'open_rv_24m', 'orig_projected_additional_accrued_interest', 'payment_plan_start_date', 'revol_bal_joint', 'sec_app_chargeoff_within_12_mths', 'sec_app_collections_12_mths_ex_med', 'sec_app_earliest_cr_line', 'sec_app_fico_range_high', 'sec_app_fico_range_low', 'sec_app_inq_last_6mths', 'sec_app_mort_acc', 'sec_app_mths_since_last_major_derog', 'sec_app_num_rev_accts', 'sec_app_open_acc', 'sec_app_open_act_il', 'sec_app_revol_util', 'settlement_amount', 'settlement_date', 'settlement_percentage', 'settlement_status', 'settlement_term', 'total_bal_il', 'total_cu_tl', 'verification_status_joint']
```

In [46]:

```
# Drop drop_feature from dataframe and check remaining features
df.drop(labels=drop_feature, axis=1, inplace=True)
df.shape # check dataframe number of rows and columns after drop some features
```

Out[46]: (176081, 86)

```
In [47]: print(df.columns)
```

```
Index(['id', 'loan_amnt', 'funded_amnt', 'funded_amnt_inv', 'term', 'int_rate',
      'installment', 'grade', 'sub_grade', 'home_ownership', 'annual_inc',
      'verification_status', 'issue_d', 'loan_status', 'pymnt_plan',
      'zip_code', 'addr_state', 'dti', 'delinq_2yrs', 'earliest_cr_line',
      'fico_range_low', 'fico_range_high', 'inq_last_6mths', 'open_acc',
      'pub_rec', 'revol_bal', 'revol_util', 'total_acc',
      'initial_list_status', 'out_prncp', 'out_prncp_inv', 'total_pymnt',
      'total_pymnt_inv', 'total_rec_prncp', 'total_rec_int',
      'total_rec_late_fee', 'recoveries', 'collection_recovery_fee',
      'last_pymnt_d', 'last_pymnt_amnt', 'last_credit_pull_d',
      'last_fico_range_high', 'last_fico_range_low',
      'collections_12_mths_ex_med', 'policy_code', 'application_type',
      'acc_now_delinq', 'tot_coll_amt', 'tot_cur_bal', 'total_rev_hi_lim',
      'acc_open_past_24mths', 'avg_cur_bal', 'bc_open_to_buy', 'bc_util',
      'chargeoff_within_12_mths', 'delinq_amnt', 'mo_sin_old_il_acct',
      'mo_sin_old_rev_tl_op', 'mo_sin_rcnt_rev_tl_op', 'mo_sin_rcnt_tl',
      'mort_acc', 'mths_since_recent_bc', 'num_accts_ever_120_pd',
      'num_actv_bc_tl', 'num_actv_rev_tl', 'num_bc_sats', 'num_bc_tl',
      'num_il_tl', 'num_op_rev_tl', 'num_rev_accts', 'num_rev_tl_bal_gt_0',
      'num_sats', 'num_tl_30dpd', 'num_tl_90g_dpd_24m', 'num_tl_op_past_12m',
      'pct_tl_nvr_dlq', 'percent_bc_gt_75', 'pub_rec_bankruptcies',
      'tax_liens', 'tot_hi_cred_lim', 'total_bal_ex_mort', 'total_bc_limit',
      'total_il_high_credit_limit', 'hardship_flag', 'disbursement_method',
      'debt_settlement_flag'],
      dtype='object')
```

```
In [48]: # further drop some features such as 'id' at Author's judgement that they're not as helpful to modeling
drop_list=['id', 'addr_state', 'funded_amnt', 'funded_amnt_inv', 'pymnt_plan', 'delinq_2yrs', 'inq_last_6mths', 'out_prncp_inv']
df.drop(labels=drop_list, axis=1, inplace=True)
```

```
In [49]: # check number of rows and columns in dataframe after drop
df.shape
```

```
Out[49]: (176081, 26)
```

```
In [50]: # define a plot function to plot features
def plot_feature (col_name, full_name, continuous):
    f, ax1=plt.subplots(1,1) # create figures
    if continuous:
        ax1=sns.boxplot(data=df, x=col_name, y='loan_status')
```

```

ax1.set_xlabel(full_name)
ax1.set_ylabel('Loan Status') # if the feature is continuous, create boxplot and set x, y labels
else:
    charge_off_rates=df.groupby(col_name)['loan_status'].value_counts(normalize=True).loc[:, 'Charged Off']
    ax1=sns.barplot(x=charge_off_rates.index, y=charge_off_rates.values)
    ax1.set_xlabel(full_name)
    ax1.set_ylabel('Charge_off_rates') # if the feature is categorical, calculate charge_off_rates grouped by feature
plt.tight_layout()


```

In [51]:

```

# choose some features to plot and explore based on business judgement
plot_feature('loan_amnt', 'Loan Amount', continuous=True)

```




In [52]:

```


# As boxplot shown above, the fully paid loan amount is lower than the charged off loan amount
plot_feature('int_rate', 'Interest Rate', continuous=True)

```




```
In [53]: # As shown above, the interquartile interest rate of Charged Off loans ranges fraom 12% to 18% while interquartile of i
#It is safe to say that Charged Off loans tend to have higher interest rate than Fully Paid loans
```

```
In [54]: plot_feature('installment', 'Installment', continuous=True)
```



```
In [55]: # We can see Charged Off loans have slightly higher installments
```

```
In [56]: plot_feature('annual_inc', 'Annual Income', continuous=True)
```




```
In [57]: # Modify annual income to log values for better visualization
df['log_annual_inc'] = df['annual_inc'].apply(lambda x: np.log10(x+1))
df['log_annual_inc'].describe()
```

```
Out[57]: count      176081.000000
mean          4.818483
std           0.236587
min           0.000000
25%           4.662767
50%           4.812920
75%           4.963793
max           6.954243
Name: log_annual_inc, dtype: float64
```

```
In [58]: # drop annual income column since we have log annual income now
df.drop(labels='annual_inc', axis=1, inplace=True)
```

```
In [59]: plot_feature('log_annual_inc', 'Log Annual Income', continuous=True)
```

```
In [60]: # As we can see, Fully Paid Loan borrowers have higher annual income than Charged Off loan borrowers
```

```
In [61]: #check grade and sub_grade
print(sorted(df['grade'].unique()))
```

```
['A', 'B', 'C', 'D', 'E', 'F', 'G']
```

```
In [62]: print(sorted(df['sub_grade'].unique()))
```


```
['A1', 'A2', 'A3', 'A4', 'A5', 'B1', 'B2', 'B3', 'B4', 'B5', 'C1', 'C2', 'C3', 'C4', 'C5', 'D1', 'D2', 'D3', 'D4', 'D5', 'E1', 'E2', 'E3', 'E4', 'E5', 'F1', 'F2', 'F3', 'F4', 'F5', 'G1', 'G2', 'G3', 'G4', 'G5']
```

```
In [63]: # As we can see, grade is included in the sub_grade, so we can drop grade
df.drop(labels='grade', axis=1, inplace=True)
```

```
In [64]: plot_feature('sub_grade', 'Subgrade', continuous=False)
plt.xticks(rotation='vertical')
```

```
Out[64]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
                17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
```

```
34 ]),  
[Text(0, 0, 'A1'),  
Text(1, 0, 'A2'),  
Text(2, 0, 'A3'),  
Text(3, 0, 'A4'),  
Text(4, 0, 'A5'),  
Text(5, 0, 'B1'),  
Text(6, 0, 'B2'),  
Text(7, 0, 'B3'),  
Text(8, 0, 'B4'),  
Text(9, 0, 'B5'),  
Text(10, 0, 'C1'),  
Text(11, 0, 'C2'),  
Text(12, 0, 'C3'),  
Text(13, 0, 'C4'),  
Text(14, 0, 'C5'),  
Text(15, 0, 'D1'),  
Text(16, 0, 'D2'),  
Text(17, 0, 'D3'),  
Text(18, 0, 'D4'),  
Text(19, 0, 'D5'),  
Text(20, 0, 'E1'),  
Text(21, 0, 'E2'),  
Text(22, 0, 'E3'),  
Text(23, 0, 'E4'),  
Text(24, 0, 'E5'),  
Text(25, 0, 'F1'),  
Text(26, 0, 'F2'),  
Text(27, 0, 'F3'),  
Text(28, 0, 'F4'),  
Text(29, 0, 'F5'),  
Text(30, 0, 'G1'),  
Text(31, 0, 'G2'),  
Text(32, 0, 'G3'),  
Text(33, 0, 'G4'),  
Text(34, 0, 'G5')] )
```




```
In [65]: # It is very clear that as sub_grade goes worse, Charge Off rates goes higher
```

```
In [66]: # check feature home_ownership
df['home_ownership'].value_counts(dropna=False)
```

```
Out[66]: MORTGAGE      86296
RENT          69952
OWN           19832
ANY              1
Name: home_ownership, dtype: int64
```

```
In [67]: plot_feature('home_ownership', 'Home Ownership', continuous=False)
```




```
In [68]: # As we can see, mortgaged home owners have lowest loan charge_off_rates, renters have the highest charge off rates. In
```

```
In [69]: # check dti, debt to income ratio  
df['dti'].describe()
```

```
Out[69]: count    176079.000000  
mean        18.911104  
std         9.180844  
min         0.000000  
25%        12.310000  
50%        18.340000  
75%        25.090000  
max         99.000000  
Name: dti, dtype: float64
```

```
In [70]: plot_feature('dti', 'Debt to Income Ratio', continuous=True)
```



```
In [71]: # Looks like there are some outliers. Remove records that have dti larger than 200
```

```
In [72]: # count how many outliers (dti>200)
(df['dti']>=40).sum()
```

Out[72]: 69

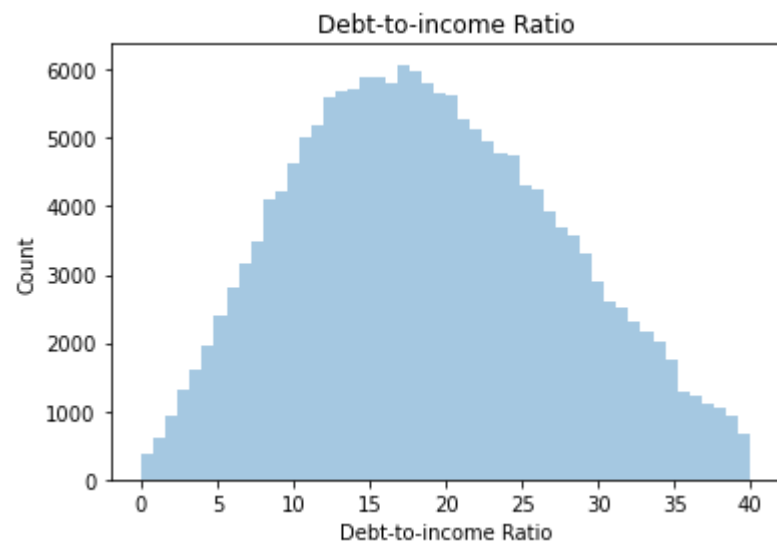
```
In [73]: # only 69 compare to total data point of 176079, 69 is very few, we can plot without these outliers
```

```
In [74]: plt.figure()
sns.distplot(df.loc[df['dti'].notnull()&(df['dti']<40), 'dti'], kde=False)
plt.xlabel('Debt-to-income Ratio')
plt.ylabel('Count')
plt.title('Debt-to-income Ratio')
```

/Users/danhongcheng/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

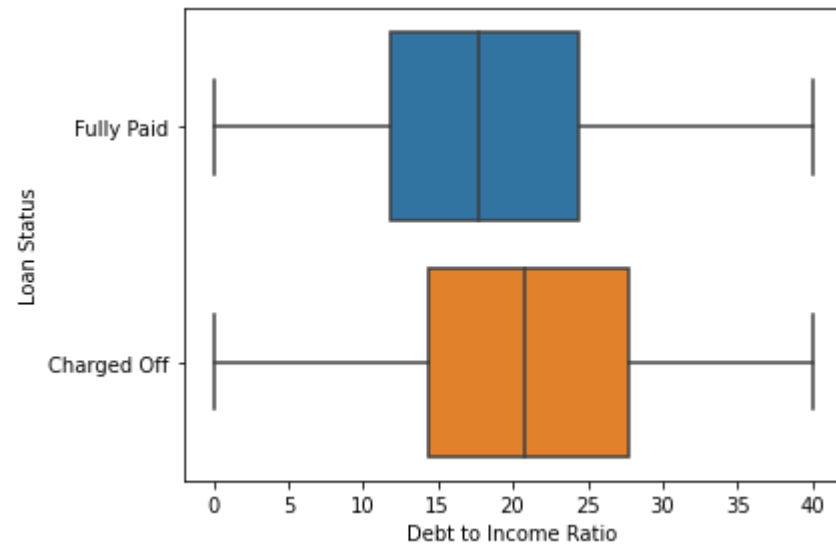
Out[74]: Text(0.5, 1.0, 'Debt-to-income Ratio')



```
In [75]: # drop rows that have dti higher than 40  
df=df[df.dti<=40]  
df['dti'].describe()
```

```
Out[75]: count    176010.00000  
mean        18.88532  
std         8.65833  
min         0.00000  
25%        12.31000  
50%        18.34000  
75%        25.08000  
max         39.99000  
Name: dti, dtype: float64
```

```
In [76]: plot_feature('dti', 'Debt to Income Ratio', continuous=True)
```



```
In [77]: # As we can see, Charged Off loan borrowers have higher debt to incomre rato than Fully Paid loan borrowers
```

```
In [78]: #Check Fico score
df[['fico_range_low', 'fico_range_high']].describe()
```

```
Out[78]:
```

	fico_range_low	fico_range_high
count	176010.000000	176010.000000
mean	694.533066	698.533197
std	31.063531	31.064167
min	660.000000	664.000000
25%	670.000000	674.000000
50%	685.000000	689.000000
75%	710.000000	714.000000
max	845.000000	850.000000


```
In [79]: # take average of fico_range_low and fico_range_high to get a fico_score
```

```
df['fico_score']=0.5*(df['fico_range_low']+df['fico_range_high'])
df['fico_score'].describe()
```

```
Out[79]: count    176010.000000
mean       696.533132
std        31.063848
min        662.000000
25%        672.000000
50%        687.000000
75%        712.000000
max        847.500000
Name: fico_score, dtype: float64
```

```
In [80]: #drop fico_range_high and fico_range_low since we have fico_score
df.drop(labels=['fico_range_high', 'fico_range_low'], axis=1, inplace=True)
```

```
In [81]: # plot fico_score
plot_feature('fico_score', 'FICO Score', continuous=True)
```



```
In [82]: # We can see Fully Paid loan borrowers have higher fico scores than Charged Off loan borrowers
```

```
In [83]: # Now as we explored some features and they all make sense when we plot them. That indicates this dataset is so far so
```



```
# First we need to change loan status to 0/1
df['charged_off']=(df['loan_status']=='Charged Off').apply(np.uint8)
print(df['charged_off'].value_counts())
```

```
0    140936
1     35074
Name: charged_off, dtype: int64
```

```
In [84]: # drop original loan_status
df.drop(labels='loan_status', axis=1, inplace=True)
```

```
In [85]: # check how many rows and features do we have
df.shape
```

```
Out[85]: (176010, 24)
```

```
In [86]: # Find columns that have categorical variables
num_cols=df._get_numeric_data().columns # find numerical columns
cols=df.columns # get all columns
cat_cols=list(set(cols)-set(num_cols)) # create a list that contain categorical columns
df=pd.get_dummies(df, columns=cat_cols, drop_first=True)
df.shape
```

```
Out[86]: (176010, 1590)
```

```
In [87]: print(df.sample(5))
```

	loan_amnt	int_rate	installment	dti	open_acc	pub_rec	revol_bal	\
173258	1000	12.69	33.55	11.10	3	0	3469	
87731	3500	9.17	111.58	0.88	11	1	1856	
136648	16000	11.53	527.85	20.21	11	0	21072	
114965	35000	18.25	1269.73	27.20	12	0	24774	
65415	5000	6.89	154.14	29.04	14	0	16197	

	revol_util	total_acc	mort_acc	...	sub_grade_F2	sub_grade_F3	\
173258	82.6	9	0	...	0	0	
87731	12.3	31	4	...	0	0	
136648	39.5	39	0	...	0	0	
114965	87.9	45	4	...	0	0	
65415	77.5	19	0	...	0	0	

	sub_grade_F4	sub_grade_F5	sub_grade_G1	sub_grade_G2	sub_grade_G3	\
173258	0	0	0	0	0	
87731	0	0	0	0	0	
136648	0	0	0	0	0	
114965	0	0	0	0	0	
65415	0	0	0	0	0	

	sub_grade_G4	sub_grade_G5	term_ 60 months
173258	0	0	0
87731	0	0	0
136648	0	0	0
114965	0	0	0
65415	0	0	0

[5 rows x 1590 columns]

```
In [110... ## Check total number of null values in dataframe
null_row_total=np.sum(df.isnull().sum())
print(null_row_total)
```

1683

```
In [111... ## Check the total number of infinity values in dataframe
infinity_row_total=np.sum(np.isinf(df).values.sum())
print(infinity_row_total)
```

0

```
In [112... ## Check total entries in the dataframe
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 176010 entries, 0 to 199998
Columns: 1590 entries, loan_amnt to term_ 60 months
dtypes: float64(7), int64(8), uint8(1575)
memory usage: 289.9 MB
```

```
In [ ]: ## As we can see, 1683 null values is small compared to total entries of 176010. We will drop these null values.
```

```
In [116... df.dropna(axis=0, inplace=True)
```

```
In [117...  
## Double check if there is still null value  
print(sum(df.isnull().sum()))
```

0

```
In [118...  
## Now we are ready to build machine learning models. We will use the following three machine learning methods:  
## 1. Random forest  
## 2. K-nearest neighbors  
## 3. Logistic regression  
## 4. Neural Network  
## And then we will compare their performances.  
  
# Create responsive variable y and predictors X  
y=df['charged_off']  
X=df.loc[:,df.columns!='charged_off']  
  
# import train_test_split from sklearn  
from sklearn.model_selection import train_test_split
```

```
In [119...  
# split the data into 80% for training and 20% for test  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, shuffle=True)
```

```
In [333...  
# 1. Random Forest Classifier  
# Import necessary modules for RandomForest Classifier
```

```
In [120...  
from sklearn.ensemble import RandomForestClassifier  
clf=RandomForestClassifier(n_estimators=100) # build randomforestclassifier clf  
clf.fit(X_train, y_train) # fit training data to the classifier  
y_pred=clf.predict(X_test) # predict charged_off using test data
```

```
In [141...  
# Evalute the performance of RandomForest Classifier  
  
from sklearn import metrics # import metrics module  
from sklearn.metrics import roc_auc_score # import roc_auc_score  
  
# Caculate accuracy, precision, recall, f1 score and roc_auc_score of the metrics  
clf_accuracy=metrics.accuracy_score(y_test, y_pred)
```

```
clf_precision=metrics.precision_score(y_test, y_pred, average='weighted')
clf_recall=metrics.recall_score(y_test, y_pred, average='weighted')
clf_f1=metrics.f1_score(y_test, y_pred, average='weighted')
clf_roc_auc=roc_auc_score(y_test, clf.predict_proba(X_test)[: , 1])
```

```
In [181... # 2. k-nearest neighbors
# import modules for k-nearest neighbors
from sklearn import neighbors
# build knn model
knn=neighbors.KNeighborsClassifier(n_neighbors=5)
# train knn model on training data
knn.fit(X_train, y_train)
```

```
Out[181... KNeighborsClassifier()
```

```
In [154... # predict charged_off using knn model and test data
y_pred_knn=knn.predict(X_test)
```

```
In [150... # evaluate performance of knn model by calculating accuracy, precision, recall, f1_score, roc_auc_score
knn_accuracy=metrics.accuracy_score(y_test, y_pred_knn)
knn_precision=metrics.precision_score(y_test, y_pred_knn, average='weighted')
knn_recall=metrics.recall_score(y_test, y_pred_knn, average='weighted')
knn_f1=metrics.recall_score(y_test, y_pred_knn, average='weighted')
knn_roc_auc=roc_auc_score(y_test, knn.predict_proba(X_test)[: ,1])
```

```
In [161... # 3. Logistic regression
# Import modules for logistic regression from sklearn
from sklearn import linear_model
# build a logistic regression classifier lrg and train on training data
lrg=linear_model.LogisticRegression(random_state=42, n_jobs=-1).fit(X_train, y_train)
y_pred_lrg=lrg.predict(X_test)
```

```
In [163... # evaluate performance of logistic regression model by calculating accuracy, prediction, recall, f1_score, roc_auc_score
lrg_accuracy=metrics.accuracy_score(y_test, y_pred_lrg)
lrg_precision=metrics.precision_score(y_test, y_pred_lrg, average='weighted')
lrg_recall=metrics.recall_score(y_test, y_pred_lrg, average='weighted')
```

```
lrg_f1=metrics.recall_score(y_test, y_pred_lrg, average='weighted')  
lrg_roc_auc=roc_auc_score(y_test, lrg.predict_proba(X_test)[: ,1])
```

In [174...

```
# 4. Neural network  
# import modules for neural network  
from sklearn.neural_network import MLPClassifier  
# import standard scaler to scale data  
from sklearn.preprocessing import StandardScaler  
scaler=StandardScaler()  
# scale train data  
scaler.fit(X_train)  
X_train=scaler.transform(X_train)  
# apply same transformation to test data  
X_test=scaler.transform(X_test)  
# build neural network multi-layer classifier mlp  
mlp=MLPClassifier(solver='lbfgs', alpha=1e-5, hidden_layer_sizes=(100, ), max_iter=1000, random_state=1)  
# train the model on training data  
mlp.fit(X_train, y_train)  
# predict charge_off  
y_pred_mlp=mlp.predict(X_test)
```

In [175...

```
# evaluate performance of neural network model by calculating accuracy, prediction, recall, f1_score, roc_auc_score  
mlp_accuracy=metrics.accuracy_score(y_test, y_pred_mlp)  
mlp_precision=metrics.precision_score(y_test, y_pred_mlp, average='weighted')  
mlp_recall=metrics.recall_score(y_test, y_pred_mlp, average='weighted')  
mlp_f1=metrics.recall_score(y_test, y_pred_mlp, average='weighted')  
mlp_roc_auc=roc_auc_score(y_test, mlp.predict_proba(X_test)[: ,1])
```

In []:

```
# build a dictionary to store models and their performance scores
```

In [201...

```
Score_list={'Accuracy':[clf_accuracy, knn_accuracy, lrg_accuracy, mlp_accuracy], 'Precision':[clf_precision, knn_precision]}
```

In [204...

```
# convert dictionary into dataframe  
Score=pd.DataFrame(Score_list, index=['Random Forest', 'K-Nearest Neighbors', 'Logistic Regression', 'Multi-layer Neural Network'])
```

In [205...

```
print(Score)
```

	Accuracy	Precision	Recall	F1	ROC_AUC
Random Forest	0.807810	0.769618	0.807810	0.748266	0.726702
K-Nearest Neighbors	0.771769	0.302817	0.771769	0.771769	0.726702
Logistic Regression	0.802191	0.742376	0.802191	0.802191	0.694921
Multi-layer Neural Network	0.720762	0.724500	0.720762	0.720762	0.609645

In []:

```
## Conclusions: 1. Random Forest has the highest accuracy at 0.807, it also has good precision, recall, F1 and ROC_AUC .  
## I would recommend random forest to classify. It can make correct predictions at the probability of 0.807 (accuracy),
```